



Replace a Cisco Nexus 9336C-FX2 shared switch

ONTAP Systems

Amanda Stroman
May 12, 2021

This PDF was generated from https://docs.netapp.com/us-en/ontap-systems/switch-cisco-9336c-shared/9336c_replace_a_cisco_nexus_9336c-fx2_shared_switch.html on May 12, 2021. Always check docs.netapp.com for the latest.

Table of Contents

Replace a Cisco Nexus 9336C-FX2 shared switch 1

Replace a Cisco Nexus 9336C-FX2 shared switch

Replace a Cisco Nexus 9336C-FX2 shared switch

Replacing a defective Nexus 9336C-FX2 shared switch is a nondisruptive procedure (NDU).

Before you begin

The following conditions must exist before performing the switch replacement in the current environment and on the replacement switch.

- Existing cluster and network infrastructure:
 - The existing cluster must be verified as completely functional, with at least one fully connected cluster switch.
 - All cluster ports must be **up**.
 - All cluster logical interfaces (LIFs) must be **up** and on their home ports.
 - The ONTAP cluster ping-cluster -node node1 command must indicate that basic connectivity and larger than PMTU communication are successful on all paths.
- Nexus 9336C-FX2 replacement switch:
 - Management network connectivity on the replacement switch must be functional.
 - Console access to the replacement switch must be in place.
 - The node connections are ports 1/1 through 1/34:
 - All Inter-Switch Link (ISL) ports must be disabled on ports 1/35 and 1/36.
 - The desired reference configuration file (RCF) and NX-OS operating system image switch must be loaded onto the switch.
 - Any previous site customizations, such as STP, SNMP, and SSH, should be copied to the new switch.

About this task

You must execute the command for migrating a cluster LIF from the node where the cluster LIF is hosted.

The examples in this procedure use the following switch and node nomenclature:

- The names of the existing Nexus 9336C-FX2 switches are *sh1* and *sh2*.
- The name of the new Nexus 9336C-FX2 switches are *newsh1* and *newsh2*.
- The node names are *node1* and *node2*.
- The cluster ports on each node are named *e3a* and *e3b*.
- The cluster LIF names are *node1_clus1* and *node1_clus2* for node1, and *node2_clus1* and *node2_clus2* for node2.
- The prompt for changes to all cluster nodes is cluster1::*>.



The following procedure is based on the following network topology:

```
cluster1::*> network port show -ipSpace Cluster
```

Node: node1

Ignore

						Speed(Mbps)	Health
Health							
Port	IPspace	Broadcast	Domain	Link	MTU	Admin/Oper	Status
Status							

e3a	Cluster	Cluster		up	9000	auto/100000	healthy
false							
e3b	Cluster	Cluster		up	9000	auto/100000	healthy
false							

Node: node2

Ignore

						Speed(Mbps)	Health
Health							
Port	IPspace	Broadcast	Domain	Link	MTU	Admin/Oper	Status
Status							

e3a	Cluster	Cluster		up	9000	auto/100000	healthy
false							
e3b	Cluster	Cluster		up	9000	auto/100000	healthy
false							

4 entries were displayed.

cluster1::*> network interface show -vserver Cluster

		Logical	Status	Network	Current	Current
Is						
Vserver	Interface	Admin/Oper	Address/Mask	Node	Port	
Home						

Cluster						
	node1_clus1	up/up	169.254.209.69/16	node1	e3a	
true						
	node1_clus2	up/up	169.254.49.125/16	node1	e3b	
true						
	node2_clus1	up/up	169.254.47.194/16	node2	e3a	
true						
	node2_clus2	up/up	169.254.19.183/16	node2	e3b	
true						

4 entries were displayed.

```
cluster1::*> network device-discovery show -protocol cdp
```

Node/	Local	Discovered		
Protocol	Port	Device (LLDP: ChassisID)	Interface	Platform

node2	/cdp			
	e3a	sh1	Eth1/2	N9K-C9336C
	e3b	sh2	Eth1/2	N9K-C9336C
node1	/cdp			
	e3a	sh1	Eth1/1	N9K-C9336C
	e3b	sh2	Eth1/1	N9K-C9336C

4 entries were displayed.

```
sh1# show cdp neighbors
```

Capability Codes: R - Router, T - Trans-Bridge, B - Source-Route-Bridge
S - Switch, H - Host, I - IGMP, r - Repeater,
V - VoIP-Phone, D - Remotely-Managed-Device,
s - Supports-STP-Dispute

Device-ID	Local Intrfce	Hldtme	Capability	Platform	Port ID
node1	Eth1/1	144	H	FAS2980	e3a
node2	Eth1/2	145	H	FAS2980	e3a
sh2	Eth1/35	176	R S I s	N9K-C9336C	Eth1/35
sh2 (FDO220329V5)	Eth1/36	176	R S I s	N9K-C9336C	Eth1/36

Total entries displayed: 4

```
sh2# show cdp neighbors
```

Capability Codes: R - Router, T - Trans-Bridge, B - Source-Route-Bridge
S - Switch, H - Host, I - IGMP, r - Repeater,
V - VoIP-Phone, D - Remotely-Managed-Device,
s - Supports-STP-Dispute

Device-ID	Local Intrfce	Hldtme	Capability	Platform	Port ID
node1	Eth1/1	139	H	FAS2980	eb
node2	Eth1/2	124	H	FAS2980	eb
sh1	Eth1/35	178	R S I s	N9K-C9336C	Eth1/35
sh1	Eth1/36	178	R S I s	N9K-C9336C	Eth1/36

Total entries displayed: 4

Steps

1. If AutoSupport is enabled on this cluster, suppress automatic case creation by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=xh
```

Where x is the duration of the maintenance window in hours.

2. Optional: Install the appropriate RCF and image on the switch, newsh2, and make any necessary site preparations.
 - a. If necessary, verify, download, and install the appropriate versions of the RCF and NX-OS software for the new switch. If you have verified that the new switch is correctly set up and does not need updates to the RCF and NX-OS software, continue to [Step 3](#).
 - b. Go to the NetApp Cluster and Management Network Switches Reference Configuration File Description Page on the NetApp Support Site.
 - c. Click the link for the Cluster Network and Management Network Compatibility Matrix, and then note the required switch software version.
 - d. Click your browser's back arrow to return to the Description page, click CONTINUE, accept the license agreement, and then go to the Download page.
 - e. Follow the steps on the Download page to download the correct RCF and NX-OS files for the version of ONTAP software you are installing.
3. On the new switch, log in as admin and shut down all the ports that will be connected to the node cluster interfaces (ports 1/1 to 1/34).

If the switch that you are replacing is not functional and is powered down, go to [Step 4](#). The LIFs on the cluster nodes should have already failed over to the other cluster port for each node.

```
newsh2# config
Enter configuration commands, one per line. End with CNTL/Z.
newsh2(config)# interface e1/1-34
newsh2(config-if-range)# shutdown
```

4. Verify that all cluster LIFs have auto-revert enabled.

```
network interface show - vservers Cluster -fields auto-revert
```

```
cluster1::> network interface show -vservers Cluster -fields auto-revert
Logical
Vserver      Interface      Auto-revert
-----
Cluster      node1_clus1    true
Cluster      node1_clus2    true
Cluster      node2_clus1    true
Cluster      node2_clus2    true
4 entries were displayed.
```

5. Verify that all the cluster LIFs can communicate:

```
cluster ping-cluster <node name>
```

```

cluster1::*> cluster ping-cluster node1
Host is node2
Getting addresses from network interface table...
Cluster node1_clus1 169.254.209.69 node1 e3a
Cluster node1_clus2 169.254.49.125 node1 e3b
Cluster node2_clus1 169.254.47.194 node2 e3a
Cluster node2_clus2 169.254.19.183 node2 e3b
Local = 169.254.47.194 169.254.19.183
Remote = 169.254.209.69 169.254.49.125
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 4 path(s)
Basic connectivity fails on 0 path(s)
.....
Detected 9000 byte MTU on 4 path(s):
Local 169.254.47.194 to Remote 169.254.209.69
Local 169.254.47.194 to Remote 169.254.49.125
Local 169.254.19.183 to Remote 169.254.209.69
Local 169.254.19.183 to Remote 169.254.49.125
Larger than PMTU communication succeeds on 4 path(s)
RPC status:
2 paths up, 0 paths down (tcp check)
2 paths up, 0 paths down (udp check)

```

6. Shut down the ISL ports 1/35 and 1/36 on the Nexus 9336C-FX2 switch sh1.

```

sh1# configure
Enter configuration commands, one per line. End with CNTL/Z.
sh1(config)# interface e1/35-36
sh1(config-if-range)# shutdown
sh1(config-if-range)#

```

7. Remove all the cables from the Nexus 9336C-FX2 sh2 switch, and then connect them to the same ports on the Nexus C9336C-FX2 newsh2 switch.
8. Bring up the ISLs ports 1/35 and 1/36 between the sh1 and newsh2 switches, and then verify the port channel operation status.

Port-Channel should indicate Po1(SU) and Member Ports should indicate Eth1/35(P) and Eth1/36(P).

This example enables ISL ports 1/35 and 1/36 and displays the port channel summary on switch sh1.

```

sh1# configure
Enter configuration commands, one per line. End with CNTL/Z.
sh1 (config)# int e1/35-36
sh1 (config-if-range)# no shutdown
sh1 (config-if-range)# show port-channel summary
Flags:  D - Down          P - Up in port-channel (members)
        I - Individual    H - Hot-standby (LACP only)
        s - Suspended     r - Module-removed
        b - BFD Session Wait
        S - Switched      R - Routed
        U - Up (port-channel)
        p - Up in delay-lacp mode (member)
        M - Not in use. Min-links not met
-----
-----
Group Port-          Type      Protocol  Member      Ports
      Channel
-----
-----
1      Po1 (SU)       Eth      LACP       Eth1/35 (P)  Eth1/36 (P)

sh1 (config-if-range)#

```

9. Verify that port e3b is up on all nodes:

`network port show ipspace Cluster`

The output should be like the following:


```
cluster1::*> network port show -ipspace Cluster
```

```
Node: node1
```

```
Ignore
```

						Speed(Mbps)	Health
Health							
Port	IPspace	Broadcast	Domain	Link	MTU	Admin/Oper	Status
Status							
-----	-----	-----	-----	-----	-----	-----	-----

e3a	Cluster	Cluster		up	9000	auto/100000	healthy
false							
e3b	Cluster	Cluster		up	9000	auto/100000	healthy
false							

```
Node: node2
```

```
Ignore
```

						Speed(Mbps)	Health
Health							
Port	IPspace	Broadcast	Domain	Link	MTU	Admin/Oper	Status
Status							
-----	-----	-----	-----	-----	-----	-----	-----

e3a	Cluster	Cluster		up	9000	auto/100000	healthy
false							
e3b	Cluster	Cluster		up	9000	auto/auto	-
false							

```
4 entries were displayed.
```

10. On the same node you used in the previous step, revert the cluster LIF associated with the port in the previous step by using the network interface revert command.

In this example, LIF node1_clus2 on node1 is successfully reverted if the Home value is true and the port is e3b.

The following commands return LIF node1_clus2 on node1 to home port e3a and displays information about the LIFs on both nodes. Bringing up the first node is successful if the Is Home column is **true** for both cluster interfaces and they show the correct port assignments, in this example e3a and e3b on node1.

```
cluster1::*> network interface show -vserver Cluster
```

Is Vserver Home	Logical Interface	Status Admin/Oper	Network Address/Mask	Current Node	Current Port
Cluster					
true	node1_clus1	up/up	169.254.209.69/16	node1	e3a
true	node1_clus2	up/up	169.254.49.125/16	node1	e3b
true	node2_clus1	up/up	169.254.47.194/16	node2	e3a
false	node2_clus2	up/up	169.254.19.183/16	node2	e3a

4 entries were displayed.

11. Display information about the nodes in a cluster:

```
cluster show
```

This example shows that the node health for node1 and node2 in this cluster is true:

```
cluster1::*> cluster show
```

Node	Health	Eligibility
node1	false	true
node2	true	true

12. Verify that all physical cluster ports are up:

```
network port show ipspace Cluster
```

```
cluster1::*> network port show -ipspace Cluster
```

Node node1

Ignore

					Speed (Mbps)	Health
Health						
Port	IPspace	Broadcast Domain	Link	MTU	Admin/Oper	Status
Status						
-----	-----	-----	-----	-----	-----	-----
e3a	Cluster	Cluster	up	9000	auto/100000	healthy
false						
e3b	Cluster	Cluster	up	9000	auto/100000	healthy
false						

Node: node2

Ignore

					Speed (Mbps)	Health
Health						
Port	IPspace	Broadcast Domain	Link	MTU	Admin/Oper	Status
Status						
-----	-----	-----	-----	-----	-----	-----
e3a	Cluster	Cluster	up	9000	auto/100000	healthy
false						
e3b	Cluster	Cluster	up	9000	auto/100000	healthy
false						

4 entries were displayed.

13. Verify that all the cluster LIFs can communicate:

`cluster ping-cluster`

```

cluster1::*> cluster ping-cluster -node node2
Host is node2
Getting addresses from network interface table...
Cluster node1_clus1 169.254.209.69 node1 e3a
Cluster node1_clus2 169.254.49.125 node1 e3b
Cluster node2_clus1 169.254.47.194 node2 e3a
Cluster node2_clus2 169.254.19.183 node2 e3b
Local = 169.254.47.194 169.254.19.183
Remote = 169.254.209.69 169.254.49.125
Cluster Vserver Id = 4294967293
Ping status:
....
Basic connectivity succeeds on 4 path(s)
Basic connectivity fails on 0 path(s)
.....
Detected 9000 byte MTU on 4 path(s):
Local 169.254.47.194 to Remote 169.254.209.69
Local 169.254.47.194 to Remote 169.254.49.125
Local 169.254.19.183 to Remote 169.254.209.69
Local 169.254.19.183 to Remote 169.254.49.125
Larger than PMTU communication succeeds on 4 path(s)
RPC status:
2 paths up, 0 paths down (tcp check)
2 paths up, 0 paths down (udp check)

```

14. Confirm the following cluster network configuration:

`network port show`

```

cluster1::*> network port show -ipspace Cluster

Node: node1

Ignore
Speed (Mbps)
Health
Port      IPspace    Broadcast Domain Link MTU    Admin/Oper    Status
Status
-----
-----
e3a      Cluster    Cluster          up   9000   auto/100000   healthy
false
e3b      Cluster    Cluster          up   9000   auto/100000   healthy
false

Node: node2

```

Ignore

		Speed (Mbps)				Health
Health	Port	IPspace	Broadcast Domain	Link MTU	Admin/Oper	Status
Status						

e3a	Cluster	Cluster	up	9000	auto/100000	healthy
false						
e3b	Cluster	Cluster	up	9000	auto/100000	healthy
false						
4 entries were displayed.						

cluster1::*> network interface show -vserver Cluster

		Logical	Status	Network	Current	Current
Is						
Vserver	Interface	Admin/Oper	Address/Mask	Node	Port	
Home						

Cluster	node1_clus1	up/up	169.254.209.69/16	node1	e3a	
true	node1_clus2	up/up	169.254.49.125/16	node1	e3b	
true	node2_clus1	up/up	169.254.47.194/16	node2	e3a	
true	node2_clus2	up/up	169.254.19.183/16	node2	e3b	
true						
4 entries were displayed.						

cluster1::> network device-discovery show -protocol cdp

Node/	Local	Discovered			
Protocol	Port	Device (LLDP: ChassisID)	Interface	Platform	

node2	/cdp				
	e3a	sh1 0/2	N9K-C9336C		
	e3b	newsh2	0/2	N9K-C9336C	
node1	/cdp				
	e3a	sh1	0/1	N9K-C9336C	
	e3b	newsh2	0/1	N9K-C9336C	
4 entries were displayed.					

sh1# show cdp neighbors

Capability Codes: R - Router, T - Trans-Bridge, B - Source-Route-Bridge
S - Switch, H - Host, I - IGMP, r - Repeater,
V - VoIP-Phone, D - Remotely-Managed-Device,
s - Supports-STP-Dispute

Device-ID ID	Local Intrfce	Hldtme	Capability	Platform	Port
node1	Eth1/1	144	H	FAS2980	e3a
node2	Eth1/2	145	H	FAS2980	e3a
newsh2 Eth1/35	Eth1/35	176	R S I s	N9K-C9336C	
newsh2 Eth1/36	Eth1/36	176	R S I s	N9K-C9336C	

Total entries displayed: 4

sh2# show cdp neighbors

Capability Codes: R - Router, T - Trans-Bridge, B - Source-Route-Bridge
S - Switch, H - Host, I - IGMP, r - Repeater,
V - VoIP-Phone, D - Remotely-Managed-Device,
s - Supports-STP-Dispute

Device-ID	Local Intrfce	Hldtme	Capability	Platform	Port ID
node1	Eth1/1	139	H	FAS2980	e3b
node2	Eth1/2	124	H	FAS2980	eb
sh1	Eth1/35	178	R S I s	N9K-C9336C	Eth1/35
sh1	Eth1/36	178	R S I s	N9K-C9336C	Eth1/36

Total entries displayed: 4

15. Enable the Ethernet switch health monitor log collection feature for collecting switch-related log files, using the following commands:

- ° `system switch ethernet log setup password`
- ° `system switch ethernet log enable-collection`

```

cluster1::*> system switch ethernet log setup-password
Enter the switch name: <return>
The switch name entered is not recognized.
Choose from the following list:
sh1
sh2
cluster1::*> system switch ethernet log setup-password
Enter the switch name: sh1
RSA key fingerprint is e5:8b:c6:dc:e2:18:18:09:36:63:d9:63:dd:03:d9:cc
Do you want to continue? {y|n}::[n] y
Enter the password: <enter switch password>
Enter the password again: <enter switch password>
cluster1::*> system switch ethernet log setup-password
Enter the switch name: sh2
RSA key fingerprint is 57:49:86:a1:b9:80:6a:61:9a:86:8e:3c:e3:b7:1f:b1
Do you want to continue? {y|n}:: [n] y
Enter the password: <enter switch password>
Enter the password again: <enter switch password>
cluster1::*> system switch ethernet log enable-collection
Do you want to enable cluster log collection for all nodes in the cluster?
y|n}: [n] y
Enabling cluster switch log collection.
cluster1::*>

```



If any of these commands return an error, contact NetApp support.

16. Move the storage ports from the old switch sh2 to the new switch newsh2.
17. Verify the storage attached to HA pair 1, shared switch newsh2 is healthy.
18. Verify the storage attached to HA pair 2, shared switch newsh2 is healthy:
`storage port show -port-type ENET`

```
storage::*> storage port show -port-type ENET
```

Node	Port	Type	Mode	Speed (Gb/s)	State	Status	VLAN ID
node1							
	e3a	ENET	storage	100	enabled	online	30
	e3b	ENET	storage	0	enabled	offline	30
	e7a	ENET	storage	0	enabled	offline	30
	e7b	ENET	storage	100	enabled	online	30
node2							
	e3a	ENET	storage	100	enabled	online	30
	e3b	ENET	storage	0	enabled	offline	30
	e7a	ENET	storage	0	enabled	offline	30
	e7b	ENET	storage	100	enabled	online	30

19. Verify that the shelves are correctly cabled:

```
storage shelf port show -fields remote- device,remote-port
```

```
cluster1::*> storage shelf port show -fields remote-device,remote-port
```

shelf	id	remote-port	remote-device
3.20	0	Ethernet1/13	sh1
3.20	1	Ethernet1/13	newsh2
3.20	2	Ethernet1/14	sh1
3.20	3	Ethernet1/14	newsh2
3.30	0	Ethernet1/15	sh1
3.30	1	Ethernet1/15	newsh2
3.30	2	Ethernet1/16	sh1
3.30	3	Ethernet1/16	newsh2

8 entries were displayed.

20. Remove the old switch sh2.

21. Repeat these steps for the switch sh1 and new switch newsh1.

22. If you suppressed automatic case creation, reenable it by invoking an AutoSupport message:

```
system node autosupport invoke -node * -type all -message MAINT=END
```


Copyright Information

Copyright © 2021 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system-without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP "AS IS" AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.